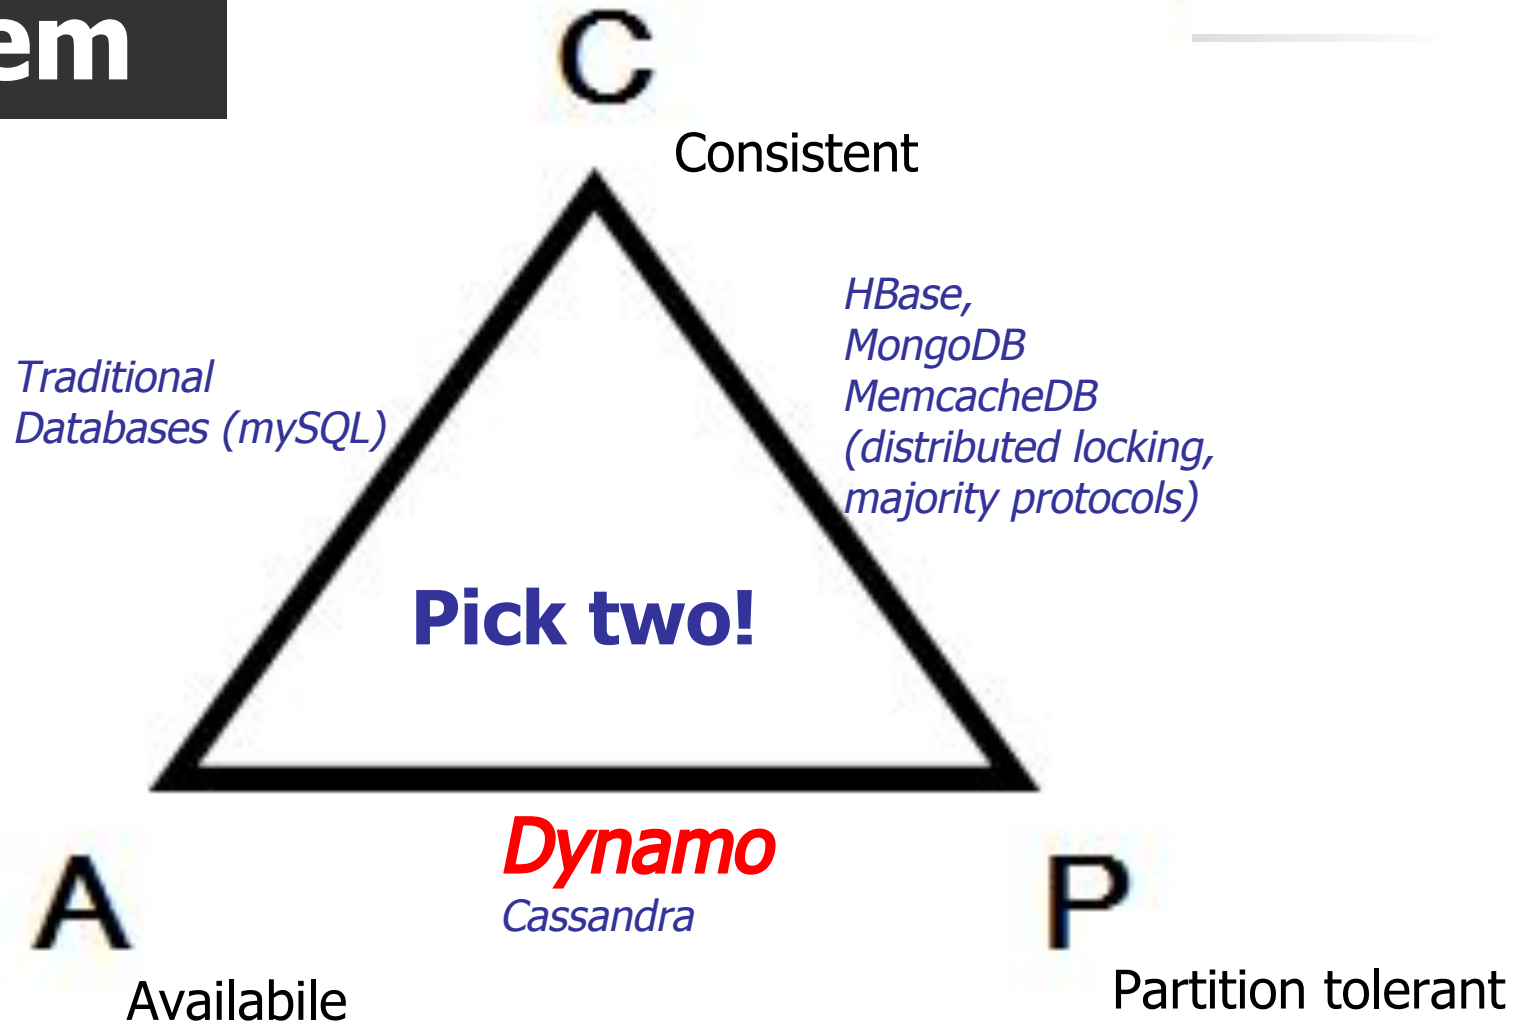



CAP Theorem



■ : Replicas

● : Clients



C+PT (-A)
(-L)

Option 1:
reads

accept reads

accept

Option 2:
~~writes~~
reads

reject writes
accept reads

reject
reject

accept writes

reject

A+PT (-C)
(+L)

writes

accept reads + writes

accept reads +

writes

‘inconsistent’ results

‘inconsistent’

results

Amazon's Dynamo:

if partitioning ☐ favor Availability (over Consistency)
normal operation ☐ favor low latency (over Consistency)

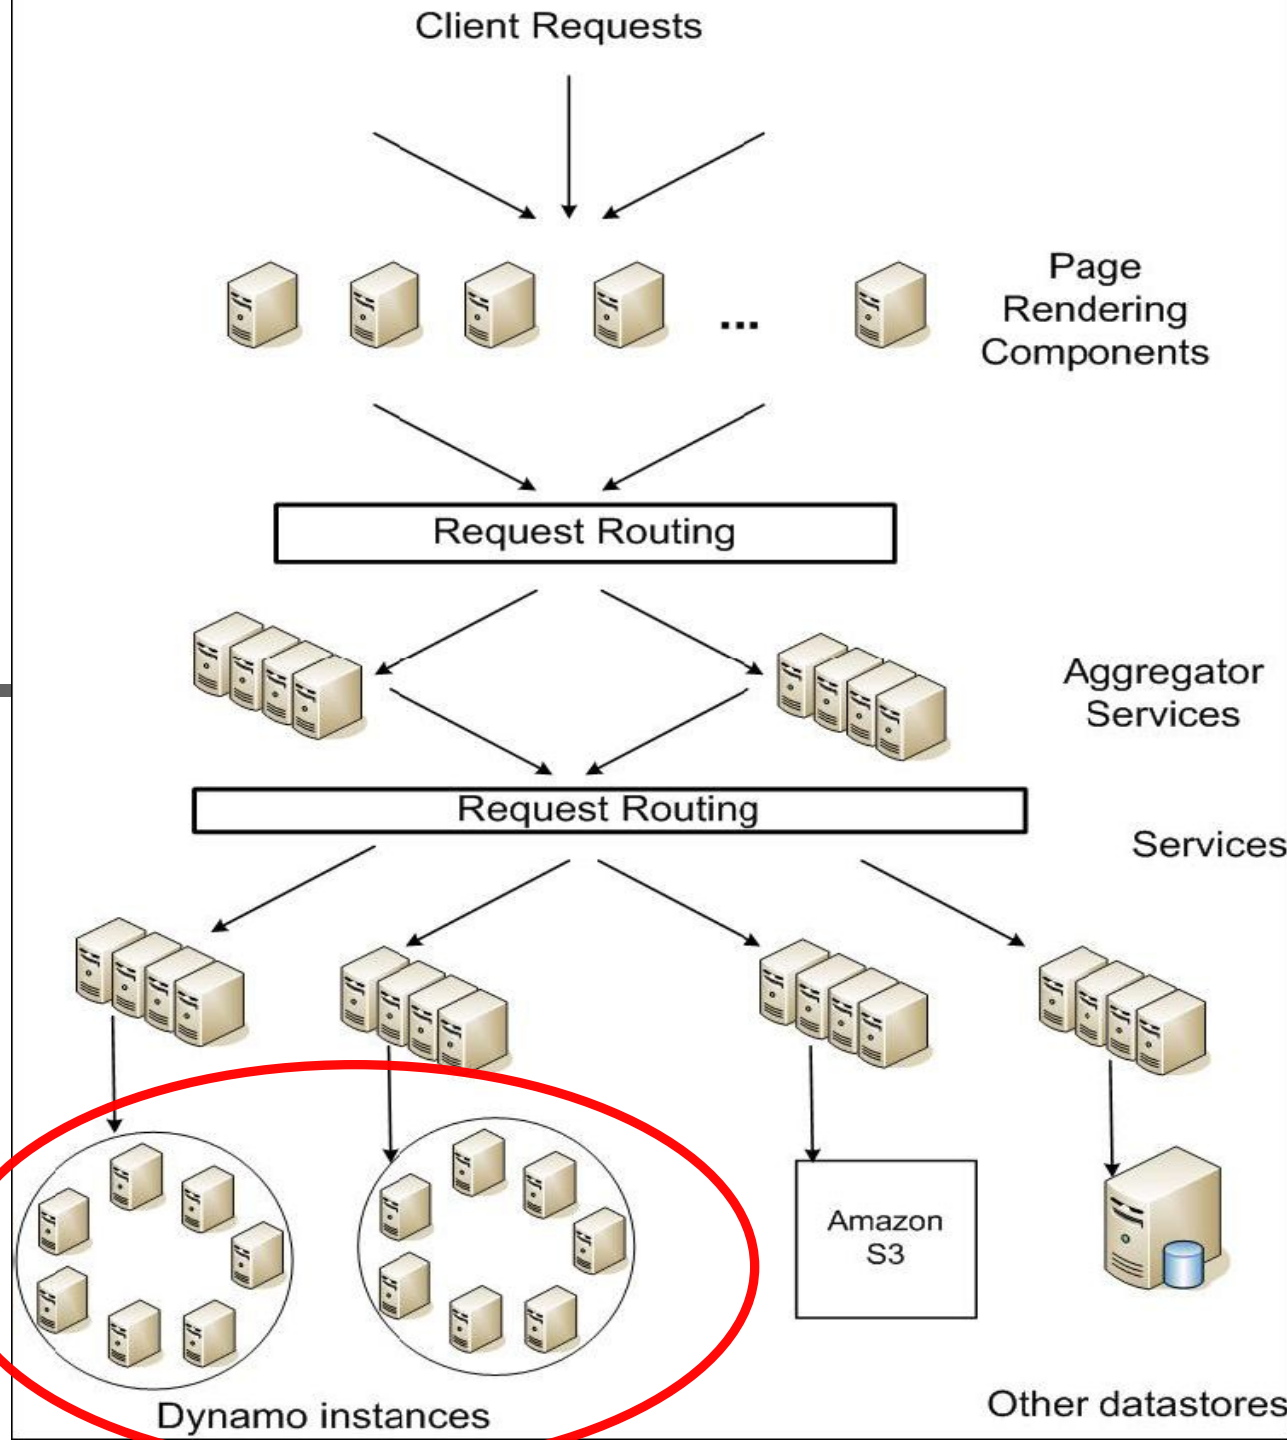
- If there is a partition (**P**):
 - How does system tradeoff A and C?
- **Else** (no partition)
 - How does system tradeoff L and C?

Amazon's Dynamo: PA/EL

<http://dbmsmusings.blogspot.com/2010/04/problems-with-cap-and-yahoos-little.html>

Dynamo: Amazon's Highly Available Key-value Store

(SOSP'07)



Giuseppe DeCandia,
Deniz Hastorun,
Madan Jampani,
Gunavardhan Kakulapati,
Avinash Lakshman, Alex
Pilchin, Swaminathan
Sivasubramanian, Peter
Vosshall
and Werner Vogels



Amazon eCommerce platform

Throughput


- 80M checkout operations per day (peak season) [2017]

Massive scale

- At this scale component-level outages are continuous!
- Slightest outage has significant financial consequences
 - \$1M/minute (2021)
 -and impacts customer trust

Application Requirements

- (1) High data *availability*; (2) always writeable data store



Why favor availability over consistency?

"even the slightest outage has significant financial consequences and impacts customer trust"

- ... consistency violations may as well have financial consequences and impact customer trust
 - But not in (a majority of) Amazon's services
 - NB: Billing is a separate story

Why favour writes vs. reads?

Architectural requirements

- Incremental scalability
- Symmetry (no 'special' node)
- Ability to run on a heterogeneous platform



Data Access Model

- Data stored as (key, object) pairs:
 - Interface *put(key, object), get(key)*
 - 'identifier' (key) generated as a hash for object
 - Objects: Opaque

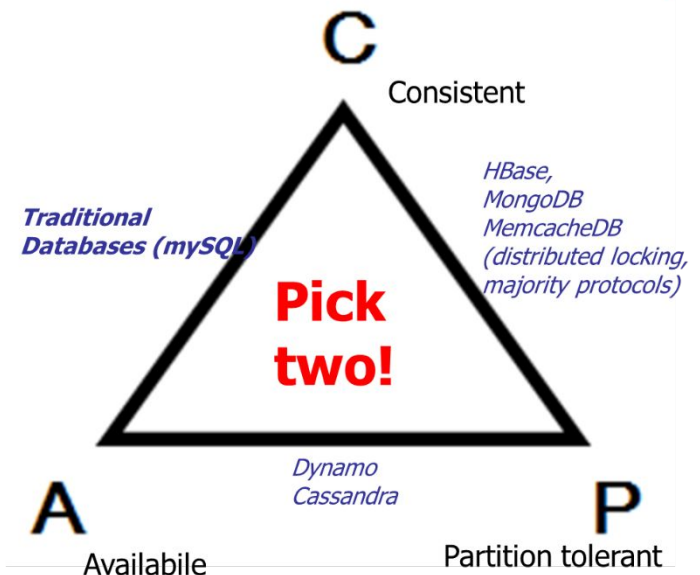
Application scenarios: shopping carts, customer preferences, session management, sales rank, product catalog, S3



Further assumptions:

- Relatively small objects (<1MB)
- Query by objectID only
- Operations do not span multiple objects
- Friendly (cooperative) environment
- One Dynamo instance per service □ 100s to 1,000s hosts/service

Why not a database?



Requirements: High availability / always writeable store / low (write) latency

Key ideas

(1) Multiple replicas ...

- ... but avoid synchronous replica coordination ...
[used by solutions that provide strong/sequential consistency]
- ... 'weak consistency' makes it possible to provide availability
- However need subsequent decisions:
 - **WHEN** to resolve possible consistency conflicts,
 - Dynamo: at read time (allows providing an "always writeable" data store)
 - **WHO** should solve them
 - Dynamo: the data store [if it can], **OR** [if that fails] the application (configurable, app specific)

(2) Vector clocks [to keep track of causality]

(3) Quorum protocols [to reduce latency / manage tradeoffs]



Key issues

- Partitioning the key/data space. Request routing
- High availability for writes
- Handling temporary node failures
- Recovering from permanent failures
- Membership and failure detection



Things to remember from last time ...

- Vector clocks

- Can be used to model causality dependence.

So far two usage examples

- [Causally ordered] Group communication
 - [Replica management] Determine replica divergence and the causally related stream of updates

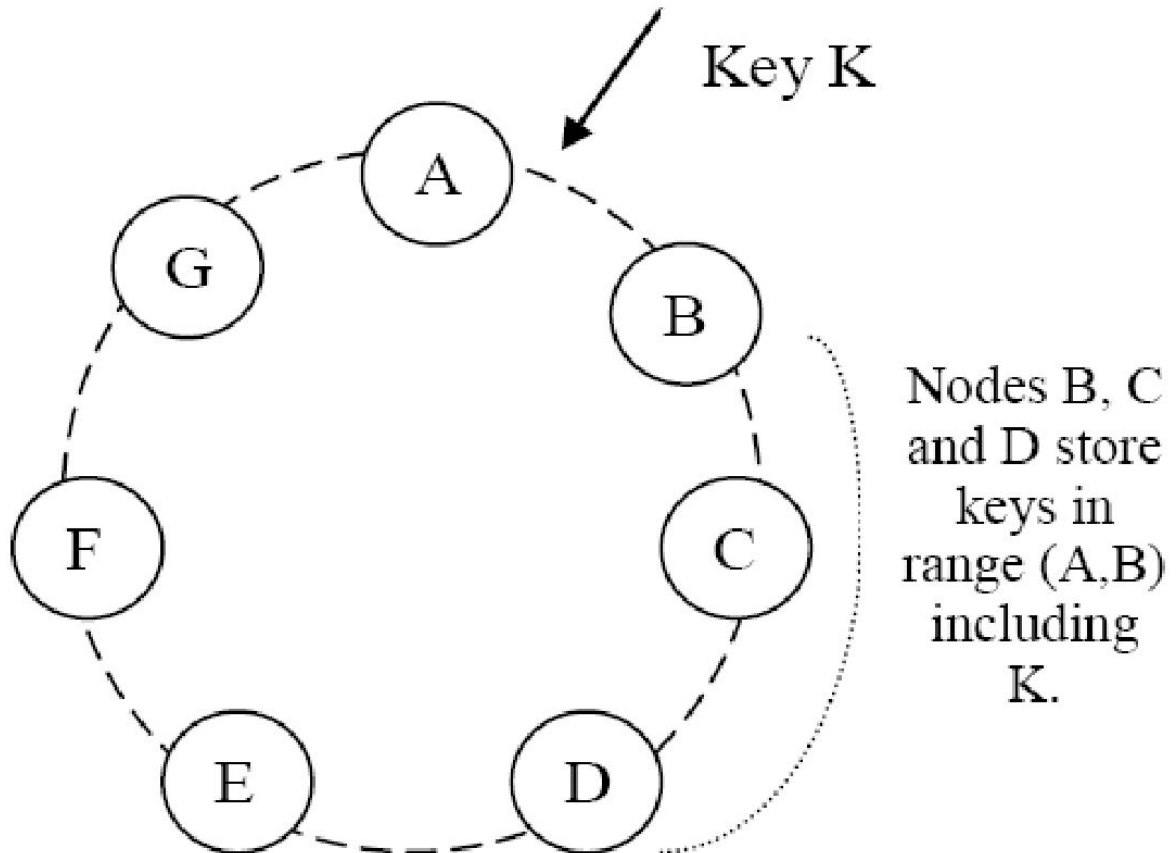
- Voting-based / Quorum-based protocols

- More examples later

Problem	Technique	Advantage
Partitioning / Request routing	<ul style="list-style-type: none"> Consistent hashing 	Incremental scalability, load balancing, etc.
High availability for writes	<ul style="list-style-type: none"> Eventual consistency Reconciliation during reads (uses vector clocks) Quorum protocol 	Low latency writes
Handling temporary failures	<ul style="list-style-type: none"> 'Sloppy' quorum protocol and hinted handoff 	Provides availability and durability when some of the replicas are temporarily not available.
Recovering from permanent failures	<ul style="list-style-type: none"> Anti-entropy using Merkle trees 	Synchronizes divergent nodes in the background.
Membership and failure detection	<ul style="list-style-type: none"> epidemic-based membership protocol 	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.

Partition Algorithm: Consistent hashing

- Each data item is replicated at N hosts (successors)



Problem	Technique	Advantage
Partitioning / Request routing	<ul style="list-style-type: none"> Consistent hashing 	Incremental scalability, load balancing, etc.
High availability for writes	<ul style="list-style-type: none"> Eventual consistency Reconciliation during reads (uses vector clocks) Quorum protocol 	Availability
Handling temporary failures	<ul style="list-style-type: none"> 'Sloppy' quorum protocol and hinted handoff 	Provides high availability and durability guarantee when some of the replicas are not available.
Recovering from permanent failures	<ul style="list-style-type: none"> Anti-entropy using Merkle trees 	Synchronizes divergent nodes in the background.
Membership and failure detection	<ul style="list-style-type: none"> epidemic-based membership protocol 	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.



Traditional quorum system:

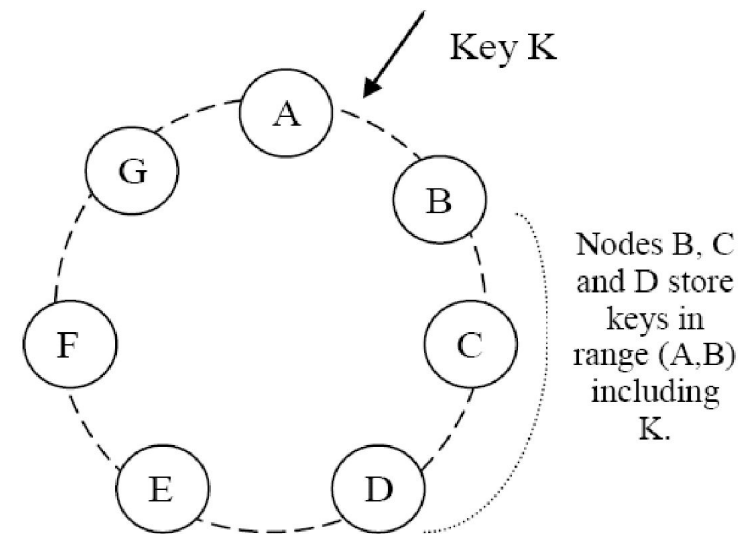
- R/W: the minimum number of nodes that must participate in a successful read/write operation.
 - Latency of a read/write operation: the slowest of the R (or W) replicas
 - To improve latency R and W are usually configured to be less than N.
- $R + W > N$ and $W > N/2$ yields a quorum-like system.

Dynamo's "Sloppy" quorum:

- A put() may return before W replicas updated!
[A first factor leading to inconsistencies]

A second factor for inconsistencies: Dealing with temporary failures through "Hinted handoff"

- Assume replication factor $R = 3$.
- Replicas on next R nodes.
- When A is temporarily down
 - (or unreachable) during a write,
 - send the write to D.
- D is 'hinted' that the replica belongs to A and will deliver it A when A recovers.
- Objective: "always writeable"





Data versioning

- Multiple replicas ...
 - ... but with focus on availability they may diverge
- The issues this introduces:
 - **when** to resolve possible conflicts?
 - Dynamo's solution: at read time (allows providing an "always writeable" data store)
 - A *put()* may return before the update has been applied at all the replicas
 - A *get()* call may return different versions of the same object.
 - **who** should solve them?
 - the data store ☐ use *vector clocks to capture causality* between different versions of the same object. **OR**
 - the application ☐ uses application specific logic.



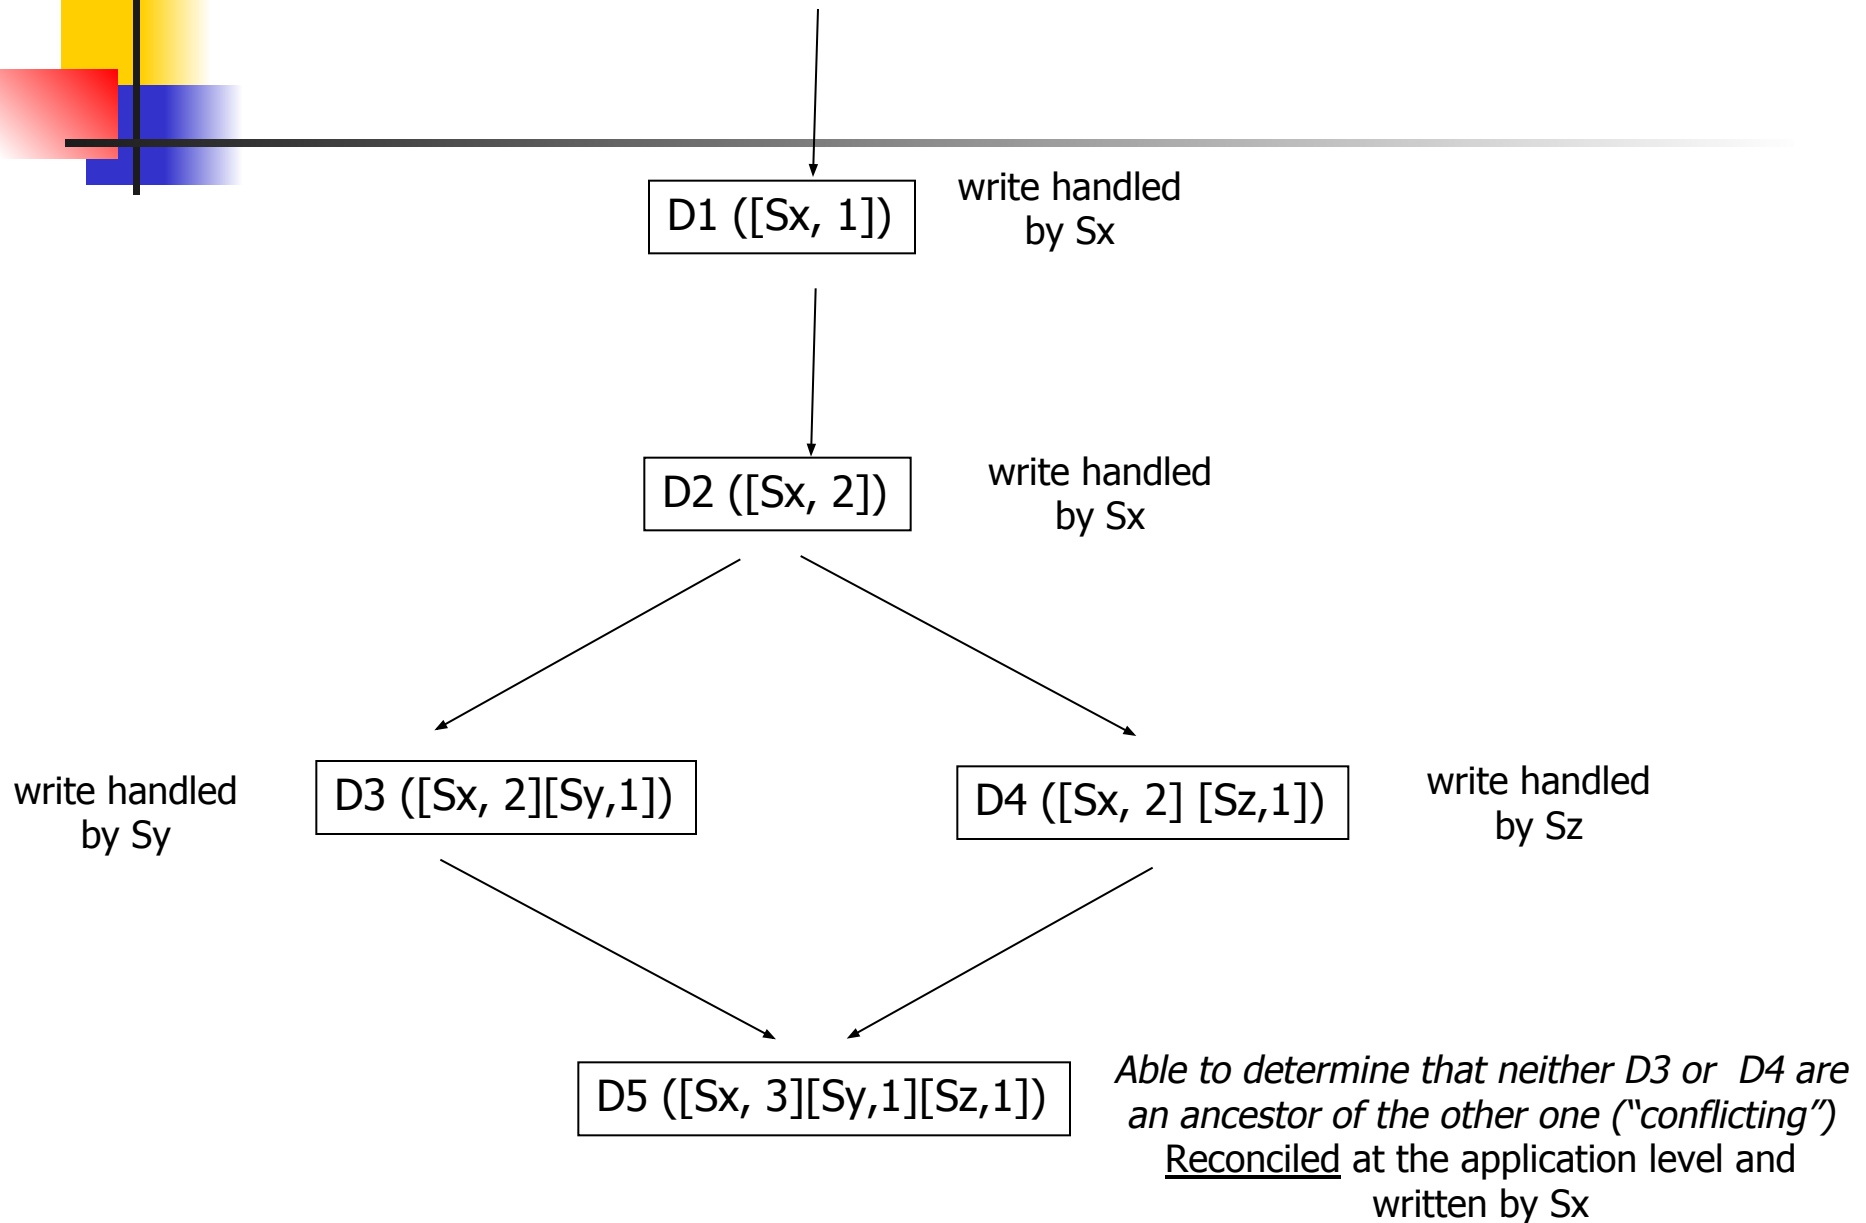

How does the data-store solve conflicts: Vector Clocks?

Each **version** of each **object** has one associated vector clock.

- list of (node, counter) pairs.

Data store: Which version to keep V' or V'' ?

- If V' is a direct ancestor of V'' then keep V''
 - direct ancestor: each counter on the V' vector clock is each less-or-equal than corresponding counter in V''
- Otherwise: application-level reconciliation





Divergent versions rarely created in practice

- 1 version ☐ 99.94%
- 2 versions ☐ 0.0057%
- 3 versions ☐ 0.00047%
- 4 versions ☐ 0.00007%

Live production environment; % versions reconciled using application logic


Source: Clients with high volume of concurrent writes
(not failures)
... these may be robots

Problem	Technique	Advantage
Partitioning	<ul style="list-style-type: none"> Consistent hashing 	Incremental scalability, load balancing, etc.
High availability for writes	<ul style="list-style-type: none"> Eventual consistency <ul style="list-style-type: none"> <i>Vector clocks with reconciliation during reads</i> Quorum protocol 	Availability
Handling temporary failures	<ul style="list-style-type: none"> 'Sloppy' quorum protocol and hinted handoff 	Provides high availability and durability guarantee when some of the replicas are not available.
Recovering from permanent failures	<ul style="list-style-type: none"> Anti-entropy using Merkle trees 	Synchronizes divergent replicas in the background.
Membership and failure detection	<ul style="list-style-type: none"> epidemic-based membership protocol 	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.

Problem	Technique	Advantage
Partitioning	<ul style="list-style-type: none"> Consistent hashing 	Incremental scalability, load balancing, etc.
High availability for writes	<ul style="list-style-type: none"> Eventual consistency <ul style="list-style-type: none"> Vector clocks with reconciliation during reads Quorum protocol 	Availability
Handling temporary failures	<ul style="list-style-type: none"> 'Sloppy' quorum protocol and hinted handoff 	Provides high availability and durability guarantee when some of the replicas are not available.
Recovering from permanent failures	<ul style="list-style-type: none"> Anti-entropy using Merkle trees 	Synchronizes divergent replicas in the background.
Membership and failure detection	<ul style="list-style-type: none"> epidemic-based membership protocol 	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.

Reconciliation

- Dynamo will replicate each data item on N successors
 - A pair (k, v) is stored by the node closest to k and replicated on N successors of that node
- Why is this may be hard?
 - As nodes may be slow, fail temporarily ...sloppy quorum, hinted handoff
 - Need to reconcile keysets





Replication Meets Epidemics

Goal: synchronize sets of key/value pairs on two nodes

■ Candidate Algorithm

- For each (k, v) stored locally, compute $\text{SHA}(k.v)$
- Every period, pick a random leaf-set neighbor
- Ask neighbor for all its hashes
- For each unrecognized hash, ask for key and value
 - (if needed use vector timestamps to reason about version freshness)


■ This is an epidemic algorithm All N members will have all (k, v) in $\log(N)$ periods

- But (above) the cost is $O(C)$, where C is the size of the set of items stored at the original node




Merkle Trees

- An efficient summarization technique
 - Interior nodes are the secure hashes of their children
 - E.g., $I = \text{SHA}(A.B)$, $N = \text{SHA}(K.L)$, etc.



Merkle Trees

- Merkle trees are an efficient summary technique
 - If the top node is signed and distributed, this signature can later be used to verify any individual block, using only $O(\log n)$ nodes, where $n = \#$ of leaves
 - E.g., to verify block C, need only R, N, I, C, & D




One use: enables client to verify integrity of data stored in a cloud


Using Merkle Trees as Summaries

- Use Merkle tree to accelerate set comparison (and identify differences)
 - B gets tree root from A, if same as local root, done
 - Otherwise, recurse down tree to find difference
[assumption: sets diverge little]

A's values:




B's values:




Using Merkle Trees as Summaries

- Use Merkle tree to accelerate set comparison (and identify differences)
 - B gets tree root from A, if same as local root, done
 - Otherwise, recurse down tree to find difference
- New cost is $O(d \log C)$
 - d = number of differences, C = size of disk

A's values:




B's values:




Using Merkle Trees as Summaries

- Still too costly:
 - If A is down for an hour, then comes back, changes will be randomly scattered throughout tree

A's values:




B's values:




Using Merkle Trees as Summaries

- Still too costly:
 - If A is down for an hour, then comes back, changes will be randomly scattered throughout tree
- Solution: order values by time instead of hash
 - Localizes values to one side of tree

A's values:



B's values:



Problem	Technique	Advantage
Partitioning	<ul style="list-style-type: none"> Consistent hashing 	Incremental scalability, load balancing, etc.
High availability for writes	<ul style="list-style-type: none"> Eventual consistency <ul style="list-style-type: none"> Vector clocks with reconciliation during reads Quorum protocol 	Availability
Handling temporary failures	<ul style="list-style-type: none"> 'Sloppy' quorum protocol and hinted handoff 	Provides high availability and durability guarantee when some of the replicas are not available.
Recovering from permanent failures	<ul style="list-style-type: none"> Anti-entropy using Merkle trees 	Synchronizes divergent replicas in the background.
Membership and failure detection	<ul style="list-style-type: none"> epidemic-based membership protocol 	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.




Dynamo Implementation


- Java
 - non-blocking IO
- Local persistence component allows for different storage engines to be plugged in:
 - Berkeley Database (BDB) Transactional Data Store: object of tens of kilobytes
 - MySQL: larger objects
- Quorum choices (N,W,R) □ influence object availability, durability, consistency

Performance evaluation

Writes



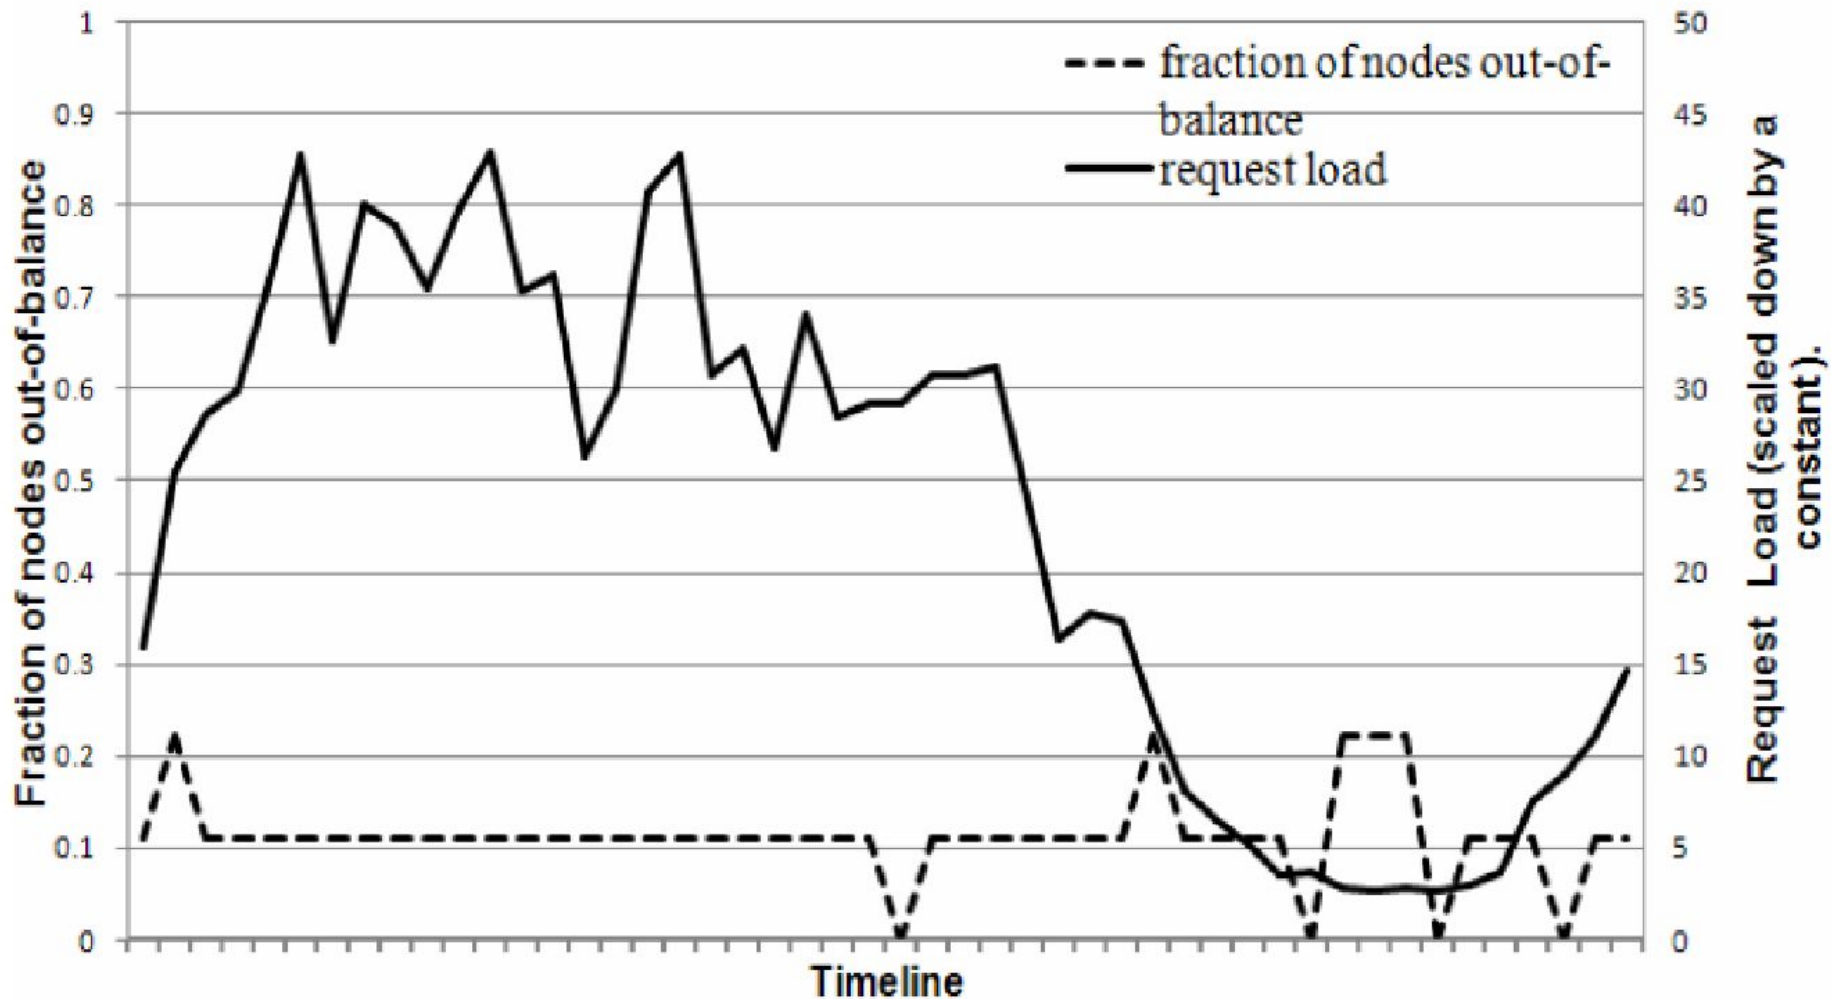
Trading between latency & durability



Comparison of performance of 99.9th %-tile latencies for buffered vs. non-buffered writes over 24 hours. The intervals between consecutive ticks in the x-axis correspond to one hour.

Load balance

out-of-balance: nodes with request load above 15% from the average system load





Divergent versions rarely created in practice

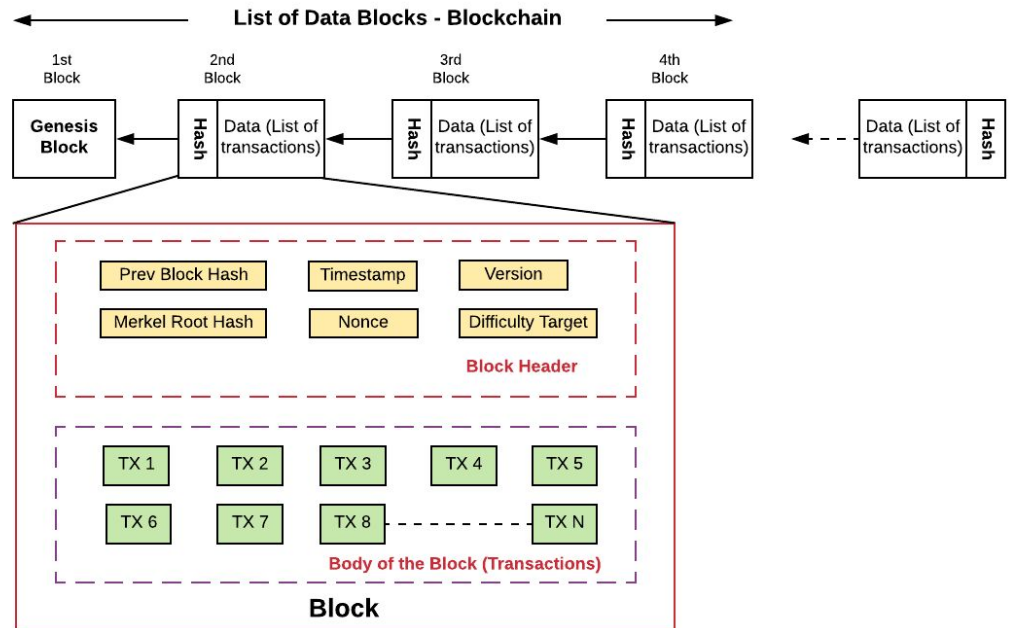
- 1 version ☐ 99.94%
- 2 versions ☐ 0.0057%
- 3 versions ☐ 0.00047%
- 4 versions ☐ 0.00007

Source: High volume of concurrent writes ... robots?

Problem	Technique	Advantage
Partitioning	Consistent Hashing	Incremental Scalability
High Availability for writes	<ul style="list-style-type: none"> Eventual consistency Vector clocks with reconciliation during reads 	Version size is decoupled from update rates.
Handling temporary failures	<ul style="list-style-type: none"> 'Sloppy' quorum and hinted handoff 	Provides high availability and durability guarantee when some of the replicas are not available.
Recovering from permanent failures	<ul style="list-style-type: none"> Anti-entropy using Merkle trees 	Synchronizes divergent replicas in the background.
Membership and failure detection	<ul style="list-style-type: none"> epidemic-based membership protocol and failure detection. 	Preserves symmetry and avoids having a centralized registry for storing membership and node liveness information.

Large set S : say 1M elements, each element 1KB-1MB

- P1: Need to test for differences between two sets
 - isDifferent (S_1, S_2) \square list of differences
- P2: You also know that if there is a difference between S_1 and S_2 , it's small
- Bonus: Verifiable proof of membership





Gilbert/Lynch theorems

Theorem 1: It is impossible in the **asynchronous** network model to implement a read/write object store that guarantees

- Availability AND
- Atomic consistency

in all executions (including those in which messages are lost)

asynchronous networks: no clocks, message delays unbounded



Gilbert/Lynch theorems

Theorem 2: It is impossible in the **partially synchronous** network model to implement a read/write object store that guarantees

- Availability AND
- Atomic consistency

in all executions (including those in which messages are lost)

partially synchronous network model. Bounds on:

- a) time it takes to deliver messages that are NOT lost, and
- b) message processing time

exist and are known, **but process clocks are not synchronized**